



TUSHAR ZAWARE

CONTACT

Wembley, HA01XY London

Mobile: +44 7448016084

tusharzaware1@gmail.com

EDUCATION

2010

Master of Science: Computer
Science

Pune University, Pune, India

CERTIFICATIONS

- SnowPro Core Certified (Snowflake)
- Aws Certified Solution Architect
- Aws Certified Developer Associate
- Microsoft Certified : AI Fundamentals
- Microsoft Certified : Azure Fundamentals
- Teradata Certified Professional
- EXIN Scrum Master
- Databricks Certified Data Engineer Associate

PROFESSIONAL SUMMARY

Senior Snowflake & Cloud Data Platform Engineer with 15+ years of experience, evolving from Teradata DBA to leading modern data platform engineering and Snowflake AI Data Cloud implementations. Proven expertise in building greenfield Snowflake platforms, designing scalable Azure infrastructure, and implementing CI/CD and Infrastructure as Code (Terraform/Bicep). Strong experience in automating RBAC/SSO, delivering secure, enterprise-grade architectures, and enabling AI-driven data solutions using Snowflake Cortex. Skilled in large-scale data migrations, performance optimisation, and cost-efficient platform design across banking, retail, and healthcare domains.

SKILLS

AI / GenAI: Snowflake Cortex AI (LLM functions, text generation, summarisation), AI-powered data applications

Infrastructure & DevOps: Terraform, Bicep, Git, Azure DevOps, GitHub Actions, CodePipeline

Programming & Scripting: Python, SQL, Shell Scripting

Monitoring & Scheduling: Apache Airflow, Control-M, Crontab

Cloud & Data Platforms: Azure, AWS, GCP, Snowflake, BigQuery, Redshift (Teradata – legacy)

Data Engineering & ETL: Matillion, Airflow, Informatica, DataStage

Database & Data Ops: Snowflake (RBAC, performance tuning, data sharing), Flyway, Liquibase, SchemaChange

BI & Visualization: Power BI (Business Objects – legacy)

WORK HISTORY

August 2022 - Current

Snowflake Data Platform Engineer, NTTDATA UK, London, United Kingdom

- Automated Snowflake resource provisioning and lifecycle management using Terraform, including end-to-end RBAC

implementation and reusable IaC modules for consistent platform provisioning.

- Extensive experience in Snowflake platform administration, including database management, architecture design, and large-scale migration projects to deliver efficient and scalable data solutions.
- Designed and maintained CI/CD pipelines using GitHub Actions and Azure DevOps for Azure, and implemented Snowflake CI/CD using native Database Change Management (DCM) along with tools like Flyway (POCs on SchemaChange and Liquibase) for controlled deployments and versioning.
- Applied DevOps best practices by integrating Infrastructure as Code (Terraform/Bicep), automated data pipeline testing, and CI/CD workflows for Data & Analytics platforms.
- Built monitoring, alerting, and observability frameworks using Snowflake-native tools, including query profiling and usage tracking for performance tuning.
- Maximised Snowflake performance and cost efficiency using clustering, materialized views, and query optimisation techniques; led cost optimisation initiatives and developed a FinOps framework for Azure and Snowflake using Streamlit.
- Implemented robust security controls including RBAC, data masking, encryption, identity federation, and network policies; performed security assessments using Azure Defender and remediated recommendations to strengthen cloud security posture.
- Established governance and compliance controls such as audit logging, access reviews, and enterprise security standards; conducted architectural reviews for Snowflake and Azure to ensure best practices and compliance.
- Designed and implemented disaster recovery and business continuity strategies for Snowflake and Azure environments.
- Enabled secure data sharing using Snowflake data sharing and reader accounts to provide external stakeholders seamless access to curated datasets.
- Implemented Azure Hub & Spoke architecture for scalable resource management, network segmentation, and enhanced security.
- Developed automated testing and validation frameworks to ensure reliability of platform changes and deployments.
- Produced high-quality documentation, including technical standards, runbooks, and onboarding guides.
- Collaborated effectively with engineering, data, and compliance teams to deliver secure, scalable, and high-performing data platforms.

2022 - Current

Snowflake Data Engineer, NTTDATA UK, London, United Kingdom

- Spearheaded comprehensive design and implementation of Snowflake-based modern data warehouse solutions within Azure environments.
- Executed greenfield setup of Snowflake platform, ensuring optimal configuration and performance.
- Developed integrated solutions with Snowflake to enhance data processing efficiency through Azure Data Factory (ADF) and Azure Data Lake Storage (ADLS Gen2).
- Engineered secure cloud-native data architectures to ensure alignment with enterprise governance and compliance standards.
- Architected and executed Snowflake Cortex service, enabling efficient semantic search functionality over extensive enterprise document datasets.
- Configured and optimized single sign-on (SSO) integration between Snowflake and Okta, leveraging SAML 2.0 for enhanced security and user experience.
- Achieved streamlined CI/CD processes for Snowflake by implementing Terraform for infrastructure and RBAC provisioning.
- Enhanced deployment efficiency through the use of Schemachange for version-controlled management of Snowflake objects.
- Administered privileged identity management (PIM) controls in Azure AD, ensuring secure and compliant access for elevated role assignments.
- Facilitated communication between data consumers and IT teams to gather complex requirements and deliver effective data solutions.
- Engineered event-driven architectures by designing and implementing streaming data pipelines in Snowflake utilizing streams, tasks, and Snowpipe.
- Facilitated integration of Snowflake and Power BI to enhance data visualization capabilities.
- Engineered a migration accelerator from BigQuery to Snowflake utilizing Python and Snowflake stored procedures.
- Leveraged Apache Airflow to orchestrate workflows, demonstrating foundational knowledge and skills in developing and supporting Directed Acyclic Graphs (DAGs). Analyzed and addressed complex data discrepancies, improving testing and validation methodologies to ensure high standards of data accuracy.

August 2022 - Current

Azure Platform Engineer, NTT DATA UK, London

- Designed and deployed an Azure cloud formation using Bicep templates, implementing a robust hub-and-spoke network architecture that supports secure, scalable resource provisioning. Ensured efficient, repeatable infrastructure-as-code (IaC) processes,

including the seamless integration of Databricks within the broader platform.

- Configured Azure Databricks clusters for optimized data processing, including custom cluster policies, security configurations, and autoscaling for cost-effective scaling.
- Established seamless integration of Databricks with Azure Data Lake Storage (ADLS) for data ingestion and processing, using access controls and managed identity authentication.
- Automated deployment pipelines via Azure DevOps, incorporating continuous integration and continuous deployment (CI/CD) for Databricks resources and Bicep code.
- Implemented Azure Privileged Identity Management (PIM) to enhance security and governance, enabling just-in-time access control and reducing the risk of excessive permissions across the cloud environment
- Integrated Azure Databricks with Power BI, enabling real-time analytics and reporting through direct query and optimized connector configurations.
- Collaborated with cross-functional teams to support analytics workloads, enhancing data flow between Databricks and Power BI for interactive data visualization.
- Implemented Databricks job scheduling, monitoring, and logging mechanisms, ensuring operational stability and prompt issue resolution across data pipelines.

August 2022 - Current

Databricks Data Engineer, NTT DATA UK, London, United Kingdom

- Designed and developed Data Factory pipelines for bulk data ingestion from relational databases into Azure Data Lake Gen2, optimizing data transfer efficiency.
- Established a scalable Azure infrastructure, configuring storage accounts, integration runtimes, service principals, app registrations, and key vaults for secure and optimized cloud utilization.
- Implemented complex data transformations using PySpark and Spark SQL in Azure Databricks, enabling efficient processing of business logic and analytics workflows.
- Extensive experience with Databricks functionalities, including Streaming, Auto Loader, Delta Live Tables, and performance tuning techniques for scalable data engineering solutions.
- Optimized Delta Lake architectures, leveraging partitioning, Z-ordering, and caching strategies to enhance query performance and large-scale data processing.
- Managed Databricks Unity Catalog, overseeing permissions, access policies, data paths, and security configurations to ensure governance, compliance, and reliable data accessibility.

- Provided technical mentorship and training, guiding teams on Databricks best practices, cost optimization, monitoring strategies, and efficient ETL/ELT implementation.
- Monitored Databricks platform usage, identifying opportunities for performance tuning, cost control, and resource optimization to enhance efficiency and scalability.

July 2021 - July 2022

Snowflake Developer, *Mistral-Data*

- Experienced in executing SAP to Snowflake migrations using Matillion ETL, automating data transfer, enhancing efficiency, and ensuring accurate data migration with robust validation and error handling.
- Collaborated with development teams to design scalable data solutions in Snowflake, gathering requirements and analyzing data structures.
- Proficient in managing Snowflake components, including multi-cluster warehouses, databases, schemas, and secure views.
- Implemented data security measures such as data masking and encryption to improve data privacy and access control.
- Demonstrated expertise in advanced Snowflake concepts like resource monitors, RBAC, query performance tuning, and data sharing features.
- Served as Matillion Administrator, managing instance setup, upgrades, and Git integration on AWS, ensuring seamless deployments and migrations.
- Organized project documentation using Confluence to promote clear communication and knowledge sharing.
- Continuously stayed current with Snowflake features and best practices, focusing on optimizing data management and system performance.

June 2015 - July 2021

Cloud Data Engineer, *Accenture*

- Skilled in Teradata-to-Snowflake migrations using Informatica, Snowpipe, and SnowSQL for ETL, and Snowflake scheduling tasks for optimized workflows.
- Proficient in developing Power BI and Excel reports from Snowflake, with expertise in query profiling, Snowflake caching, warehouse capacity management, and clustering for performance tuning.
- Experienced in managing AWS environments, including configuring and maintaining AWS services, IAM, and security for seamless cloud operations.

March 2013 - June 2015

Teradata DBA, *Teradata India Pvt. Ltd*

- Managed a 34TB Teradata production environment with 24×7 support responsibilities, covering system monitoring, performance tuning (DBQL, statistics, PDCR), spaceTeradata object administration, space management, backup & recovery, and coordination of P1 incidents/outages.

December 2010 - March 2013

Software Engineer, *Wipro Technologies*

- Responsible for Teradata and ETL development, as well as production support, utilizing Agile methodology for requirements analysis, system design, testing, and implementation.
- Developed ETL processes using Informatica Powercenter, and monitored database activity through Teradata Viewpoint and Informatica job monitoring.